



Copyright 2019. Created by Jose Marcial Portilla.

## PyTorch Gradients

This section covers the PyTorch [autograd](https://pytorch.org/docs/stable/autograd.html) implementation of gradient descent. Tools include:

- [torch.autograd.backward\(\)](https://pytorch.org/docs/stable/autograd.html#torch.autograd.backward)
- [torch.autograd.grad\(\)](https://pytorch.org/docs/stable/autograd.html#torch.autograd.grad)

Before continuing in this section, be sure to watch the theory lectures to understand the following concepts:

- Error functions (step and sigmoid)
- One-hot encoding
- Maximum likelihood
- Cross entropy (including multi-class cross entropy)
- Back propagation (backprop)

### Additional Resources:

[PyTorch Notes: \(https://pytorch.org/docs/stable/notes/autograd.html\)](https://pytorch.org/docs/stable/notes/autograd.html) Autograd mechanics

## Autograd - Automatic Differentiation

In previous sections we created tensors and performed a variety of operations on them, but we did nothing to store the sequence of operations, or to apply the derivative of a completed function.

In this section we'll introduce the concept of the *dynamic computational graph* which is comprised of all the *Tensor* objects in the network, as well as the *Functions* used to create them. Note that only the input Tensors we create ourselves will not have associated Function objects.

The PyTorch [autograd](https://pytorch.org/docs/stable/autograd.html) package provides automatic differentiation for all operations on Tensors. This is because operations become attributes of the tensors themselves. When a Tensor's `.requires_grad` attribute is set to True, it starts to track all operations on it. When an operation finishes you can call `.backward()` and have all the gradients computed automatically. The gradient for a tensor will be accumulated into its `.grad` attribute.

Let's see this in practice.

## Back-propagation on one step

We'll start by applying a single polynomial function  $y = f(x)$  to tensor  $x$ . Then we'll backprop and print the gradient  $\frac{dy}{dx}$ .

$$\text{Function : } y = 2x^4 + x^3 + 3x^2 + 5x + 1$$

$$\text{Derivative : } y' = 8x^3 + 3x^2 + 6x + 5$$

### Step 1. Perform standard imports

In [1]:

```
import torch
```

### Step 2. Create a tensor with requires\_grad set to True

This sets up computational tracking on the tensor.

In [2]:

```
x = torch.tensor(2.0, requires_grad=True)
```

### Step 3. Define a function

In [3]:

```
y = 2*x**4 + x**3 + 3*x**2 + 5*x + 1  
print(y)
```

```
tensor(63., grad_fn=<AddBackward0>)
```

Since  $y$  was created as a result of an operation, it has an associated gradient function accessible as `y.grad_fn`

The calculation of  $y$  is done as:

$$y = 2(2)^4 + (2)^3 + 3(2)^2 + 5(2) + 1 = 32 + 8 + 12 + 10 + 1 = 63$$

This is the value of  $y$  when  $x = 2$ .

### Step 4. Backprop

In [4]:

```
y.backward()
```

### Step 5. Display the resulting gradient

In [5]:

```
print(x.grad)  
tensor(93.)
```

Note that `x.grad` is an attribute of tensor  $x$ , so we don't use parentheses. The computation is the result of

$$y' = 8(2)^3 + 3(2)^2 + 6(2) + 5 = 64 + 12 + 12 + 5 = 93$$

This is the slope of the polynomial at the point (2, 63).

## Back-propagation on multiple steps

Now let's do something more complex, involving layers  $y$  and  $z$  between  $x$  and our output layer  $out$ .

### 1. Create a tensor

In [6]:

```
x = torch.tensor([[1.,2,3],[3,2,1]], requires_grad=True)
print(x)

tensor([[1., 2., 3.],
        [3., 2., 1.]], requires_grad=True)
```

### 2. Create the first layer with $y = 3x + 2$

In [7]:

```
y = 3*x + 2
print(y)

tensor([[ 5.,  8., 11.],
        [11.,  8.,  5.]], grad_fn=<AddBackward0>)
```

### 3. Create the second layer with $z = 2y^2$

In [8]:

```
z = 2*y**2
print(z)

tensor([[ 50., 128., 242.],
        [242., 128.,  50.]], grad_fn=<MulBackward0>)
```

### 4. Set the output to be the matrix mean

In [9]:

```
out = z.mean()
print(out)

tensor(140., grad_fn=<MeanBackward1>)
```

### 5. Now perform back-propagation to find the gradient of $x$ w.r.t $out$

(If you haven't seen it before, w.r.t. is an abbreviation of *with respect to*)

In [10]:

```
out.backward()
print(x.grad)

tensor([[10., 16., 22.],
        [22., 16., 10.]])
```

You should see a 2x3 matrix. If we call the final out tensor " $o$ ", we can calculate the partial derivative of  $o$  with respect to  $x_i$  as follows:

$$o = \frac{1}{6} \sum_{i=1}^6 z_i$$

$$z_i = 2(y_i)^2 = 2(3x_i + 2)^2$$

To solve the derivative of  $z_i$  we use the [chain rule](https://en.wikipedia.org/wiki/Chain_rule) ([https://en.wikipedia.org/wiki/Chain\\_rule](https://en.wikipedia.org/wiki/Chain_rule)), where the derivative of  $f(g(x)) = f'(g(x))g'(x)$

In this case

$$\begin{aligned} f(g(x)) &= 2(g(x))^2, & f'(g(x)) &= 4g(x) \\ g(x) &= 3x + 2, & g'(x) &= 3 \\ \frac{dz}{dx} &= 4g(x) \times 3 = 12(3x + 2) \end{aligned}$$

Therefore,

$$\frac{\partial o}{\partial x_i} = \frac{1}{6} \times 12(3x + 2)$$

$$\left. \frac{\partial o}{\partial x_i} \right|_{x_i=1} = 2(3(1) + 2) = 10$$

$$\left. \frac{\partial o}{\partial x_i} \right|_{x_i=2} = 2(3(2) + 2) = 16$$

$$\left. \frac{\partial o}{\partial x_i} \right|_{x_i=3} = 2(3(3) + 2) = 22$$

## Turn off tracking

There may be times when we don't want or need to track the computational history.

You can reset a tensor's `requires_grad` attribute in-place using `.requires_grad_(True)` (or `False`) as needed.

When performing evaluations, it's often helpful to wrap a set of operations in with `torch.no_grad()`:

A less-used method is to run `.detach()` on a tensor to prevent future computations from being tracked. This can be handy when cloning a tensor.

**A NOTE ABOUT TENSORS AND VARIABLES:** Prior to PyTorch v0.4.0 (April 2018) Tensors (`torch.Tensor`) only held data, and tracking history was reserved for the Variable wrapper (`torch.autograd.Variable`). Since v0.4.0 tensors and variables have merged, and tracking functionality is now available through the `requires_grad=True` flag.

